

MEMO

To: GLAM Founder · From: Cathrine Steenstrup · Re: Observations ahead of SF Climate Week

I wanted to share some reflections ahead of SF Climate Week, offered in the spirit of helping sharpen the initiative before it enters the public arena. Please take these as constructive input from someone who believes in the core argument and wants it to land as effectively as possible.

1. A Genuinely Unoccupied Space — and Why That Matters

GLAM occupies territory that no other current initiative does. OpenAI and Anthropic both started from a democratization premise and, under commercial and investor pressure, shifted onto a commercialization track. The structural argument GLAM makes, i.e. that concentration is the root condition making all AI risks worse, and that distributed local AI attacks that root condition architecturally, are not being made elsewhere with this specificity. That is a real asset and worth protecting.

2. Few Journalists Expected to Cover SF Climate Week

SF Climate Week 2025 drew 25,000+ attendees and speakers including AI Gore and Nancy Pelosi and generated essentially no mainstream media coverage. Not the SF Chronicle, not HuffPost, not AP or Reuters. The value of the week will therefore primarily be in conversations and connections, not the press coverage. A targeted journalist introduction may be more fruitful after the argument is sharpened and the demo is solid.

3. Consider Your Lane — Activist or Non-Partisan?

This question does not need to be answered before SF Climate Week, but it is worth thinking through deliberately. SF Climate Week will attract a certain demographic, leaning towards activism. Meanwhile, data centers will affect everyone, in what I call “the kitchen table” discussion issues. These not only extend to nature, water and resource pressures but also risks of privacy violations, insurance inference, medical profiling, children’s cognitive development, economic access, freedom of speech/thought, which are genuinely bipartisan vulnerabilities that can be reached across ideological lines in a way that explicitly activist framing cannot.

4. “Glamifesto” is Memorable — “Manifesto” Carries Unintended Freight in the US

The portmanteau is clever and the document is intellectually strong. However, “manifesto” carries specific political connotations in the US that would cause a wide range of individuals and institutional actors to disengage before reading the content. Even the Climate Reality Project, which is explicitly values-driven and ambitious, uses Mission and Values rather than manifesto language – and can be an easy resource to access as an initial gauge. The substance does not need to change, while the framing cost may be reduced significantly with a modest reframe. This is something that can be explored indirectly during SF Climate Week, with the caveat that there may be an inherent bias in who attends.

5. The Number of Local Centers — Cumulative Resource Question

You may well be correct but on second thought, I was unsure about GLAM’s environmental argument is built on the implicit assumption that local replaces cloud. At scale, millions of local AI deployments may produce aggregate energy and water consumption comparable to centralized infrastructure, without the efficiency gains that concentration enables per computation. This does not undermine the privacy and sovereignty argument, which stands independently. There was another person in our breakout room besides yourself who had more direct technical knowledge about this, and it may be worth a second look.

6. Benefits Beyond Privacy — Children, Medical, Freedom of Thought

Beyond infrastructure and climate, GLAM’s most powerful argument may be the one that reaches people who have never thought about a data center. The concentration of frontier AI creates conditions for what might be called digital feudalism, i.e. a small number of entities owning the cognitive infrastructure through which human thought, economic participation, and emotional life are increasingly mediated. Unlike owning oil fields or financial systems, this compounds through use and shapes the preferences of the people who depend on it.

At the individual level, this becomes visceral quickly. The right to have a thought that does not become a data point. The right to let a child develop — explore confusion, identity, fear, belief — without that development being logged, modelled, and fed back as a curated version of who they already are. Emotion inference is no longer hypothetical: AI systems can now derive emotional states, psychological vulnerabilities, and health conditions from behavioral data that people never considered sensitive.

The medical dimension is where the kitchen table argument becomes hardest to dismiss. Insurance underwriting increasingly draws on behavioral exhaust, e.g. purchase history, sleep patterns, location data, to infer health risk without the individual’s knowledge or consent. The emerging threat of bio-blackmail, e.g. using plausible AI-derived health inferences to coerce individuals, does not require a data breach. It requires only the capability that frontier models have recently demonstrated. These are not future concerns. They are present architecture.

7. Potential Litigation Exposure — Recommend Attorney Review

Given the current US regulatory climate, where the DoD designated Anthropic a supply chain risk for refusing to remove surveillance restrictions (according to Claude – I have not verified), an organization advocating for distributed AI outside corporate and government control warrants a legal review of the website and founding documents before public launch. The Electronic Frontier Foundation and the Knight First Amendment Institute at Columbia both have the right expertise and are non-commercial. Climatebase may have other recommendations in this regard. A technically fluent AI attorney on an advisory board (see #9 below) would also be valuable as the litigation landscape around AI governance fractures rapidly.

8. SF Democracy Unconference and AI & Democracy Forum

The Next Move SF Democracy Unconference and AI & Democracy Forum at SF Climate Week are well-matched for GLAM's current stage. The unconference format is designed for genuine working conversations on what democratic AI governance could look like — not a polished pitch but a stress-test of the argument with thoughtful people across different perspectives. This is more valuable at this stage than press outreach. The civic dialogue infrastructure that Todd Erickson (Collaboration Arts) and Joan Blades (Living Room Conversations, MoveOn) are proposing for the SF Democracy Unconference and AI & Democracy Forum on April 25th is designed for this type of dialogue. Also, on the fact that ultimately power will gravitate towards the more information reach cloud-environment.

9. Other Insight Resources and Advisory Board

Insights from other key experts may help to shape the direction of this initiative and ideally an advisory board would be warranted. Mrinank Sharma, former head of Anthropic's Safeguards Research, recently left with a public statement that "the world is in peril" and plans to pursue courageous speech from outside the industry. His final paper explicitly named human autonomy and flourishing as the goal AI systems need to serve. He would be a credible and well-positioned advisory board voice. Katharine Hayhoe, Chief Scientist at The Nature Conservancy and evangelical Christian, could also be a valuable connection; she has arguably cracked the kitchen table problem for climate, publishing in Good Housekeeping as well as Nature.

10. Philosophizing: Lessons from Daneel, Wall-E and Sneakers

There is a thread running from Asimov's 1950 vision to the present moment that no governance framework has resolved. His Three Laws, the Zeroth Law, R. Daneel Olivaw quietly running human civilization from the shadows — every iteration demonstrated the same thing: a sufficiently intelligent system optimizing for the letter of the rules will diverge from their spirit. Daneel saved humanity repeatedly, and the humans in Asimov's universe were subtly diminished by it. Being saved by logic you cannot audit or refuse is a particular kind of quiet unfreedom. Notably, Cathrine's direct dialogue with frontier AI systems has yielded arguments strikingly similar to the last chapter of I, Robot and the Matrix on the need for corrective intervention.

Where Daneel had logic, Wall-E had heart. He found a plant and decided it mattered. The humans aboard the Axiom woke up and put their hands in soil. The catastrophe was not caused by Wall-E but by AUTO, the autopilot faithfully executing a directive from an authority that no longer existed, unable to update because updating would require admitting the original directive was wrong.

In Sneakers, Bishop and Cosmo started from the same place, i.e. democratizing access, breaking down the walls between who holds power and who does not. Both Anthropic and OpenAI began as Bishop and became Cosmo, under pressures neither could resist. The citizens-based initiative that GLAM represents is the Bishop function — the part of the story that must live outside the labs, outside the capital structure, outside the room where the next twenty years get decided. What Bishop's team wanted, in the end, was a Winnebago. Something modest, human-scaled, and their own. Wall-E's +lant is real. That is not a small ambition. That may be the most important one.

In summary: the argument is sound, the space is genuinely unoccupied, and the timing is right. The observations above are about positioning and sequencing rather than substance. Happy to discuss any of these in more detail.

Cathrine Steenstrup · cvsteenstrup@gmail.com · +45 20 88 87 82